

## CS2102 Project (Part 2)

The objective of this part of the team project is for you to apply what you have learned in class to implement triggers and routines for a schema designed for the application mentioned in Part 1. In relation to this, we provide a reference ER diagram in “ER.pdf”, as well as a relation schema in “schema.sql”. Note that the schema has implemented some constraints using foreign keys and checks. You are to implement additional triggers and routines, **using PostgreSQL 14**, according to the requirements detailed in the following.

Note: Please **do not** make any changes to the schema in “schema.sql”. This is because we will evaluate each team’s implementation using automated testing, for which we need each team’s schema to be exactly the same. (For details, you may refer to the FAQ in Section 4.)

### 1. Constraints to be Enforced using Triggers

Please implement appropriate triggers to enforce the following 14 constraints. For simplicity, you only need to consider INSERT triggers (i.e., triggers that are activated by insertions), and do NOT need to consider DELETE or UPDATE. You can assume that you are given a database in a correct state. For each constraint, you are allowed to use multiple triggers to enforce it.

Delivery\_requests related:

- (1) Each delivery request has at least one package.

Package related:

- (2) For each delivery request, the IDs of the packages should be consecutive integers starting from 1.

Unsuccessful\_pickups related:

- (3) For each delivery request, the IDs of the unsuccessful pickups should be consecutive integers starting from 1.
- (4) The timestamp of the first unsuccessful pickup should be after the submission\_time of the corresponding delivery request. In addition, each unsuccessful pickup’s timestamp should be after the previous unsuccessful pickup’s timestamp (if any).

Legs related:

- (1) For each delivery request, the IDs of the legs should be consecutive integers starting from 1.
- (2) For each delivery request, the start time of the first leg should be after the submission\_time of the delivery request and the timestamp of the last unsuccessful pickup (if any).
- (3) For each delivery request, a new leg cannot be inserted if its start\_time is before the end\_time of the previous leg, or if the end\_time of the previous leg is NULL.

Unsuccessful\_deliveries:

- (4) The timestamp of each unsuccessful\_delivery should be after the start\_time of the corresponding leg.
- (5) For each delivery request, there can be at most three unsuccessful\_deliveries.

Cancelled\_requests related:

- (6) The cancel\_time of a cancelled request should be after the submission\_time of the corresponding delivery request.

Return\_legs related:

- (7) For each delivery request, the first return\_leg's ID should be 1, the second return\_leg's ID should be 2, and so on.
- (8) For a delivery request, the first return\_leg cannot be inserted if (i) there is no existing leg for the delivery request or (ii) the last existing leg's end\_time is after the start\_time of the return\_leg. In addition, the return\_leg's start\_time should be after the cancel\_time of the request (if any).
- (9) For each delivery request, there can be at most three unsuccessful\_return\_deliveries.

Unsuccessful\_return\_deliveries related:

- (10) The timestamp of each unsuccessful\_return\_delivery should be after the start\_time of the corresponding return\_leg.

## 2. Routines

Routines must be implemented with the same name specified in this document, following the same order of input parameters and returning the exact output parameter type (if any). You may change the names of the input parameters in your implemented routines, but not the data type of the input parameters. If a call to a routine should fail because of database constraint violations, the routine should not return any values. You may raise an exception explicitly, or simply allow it to fail silently. It is important to adhere to these instructions since automated testing will be used to evaluate your implementation. Therefore, the input and output formats of routines have to be fixed.

For all routines, you may assume that the input values are all valid, i.e., you do not need to do validity checks for them in your implementation.

In addition, for arrays, you should follow PostgreSQL's default behaviour of starting the index subscript from 1. You can also use any functions available in PostgreSQL.

### 2.1 Procedures

The following routines in this section do not have return values, and should be implemented as PostgreSQL procedures (<https://www.postgresql.org/docs/14/sql-createprocedure.html>).

- (1) submit\_request( customer\_id INTEGER, evaluator\_id INTEGER, pickup\_addr TEXT, pickup\_postal TEXT, recipient\_name TEXT, recipient\_addr TEXT, recipient\_postal TEXT, submission\_time TIMESTAMP, package\_num INTEGER, reported\_height INTEGER[], reported\_width INTEGER[], reported\_depth INTEGER[], reported\_weight INTEGER[], content TEXT[], estimated\_value NUMERIC[] )
  - Submit a delivery request from a customer
  - package\_num refers to the number of packages pertinent to the delivery request
  - The status of the delivery request should be set to "submitted"
  - The pickup\_date, num\_days\_needed, and price of the delivery request should be set to NULL
  - NUMERIC[] refers to an array of NUMERIC; TEXT[] refers to an array of TEXT
  - reported\_height[i] refers to the reported\_height of the i-th package of the delivery request; similarly for reported\_width[i], reported\_depth[i], reported\_weight[i], content[i], and estimated\_value[i]

- the actual\_height, actual\_width, actual\_depth, actual\_weight of each package should be NULL
- (2) resubmit\_request( request\_id INTEGER, evaluator\_id INTEGER, submission\_time TIMESTAMP, reported\_height INTEGER[], reported\_width INTEGER[], reported\_depth INTEGER[], reported\_weight INTEGER[] )
- Resubmit a delivery request as a new request when the actual measurements of the packages are found to be significantly different from the reported measurements during a delivery pickup.
  - The attributes of the new request are the same as those of the original one, except that:
    - the id of the new request is new
    - the evaluator\_id of the new request is set to the evaluator\_id in the input parameters
    - the status of the new request is set to “submitted”
    - the submission\_time of the new request is set to the submission\_time in the input parameters of the procedure.
    - the pickup\_date, num\_days\_needed, and price of the new request are set to NULL
  - The packages of the new request are the same as the packages in the original request, except that:
    - the request\_id of each package is the id of the new request
    - the reported\_height of the i-th package equals reported\_height[i] in the input parameters; similarly for reported\_width[i], reported\_depth[i], and reported\_weight[i]
    - the actual\_height, actual\_width, actual\_depth, and actual\_weight of each package are set to NULL
- (3) insert\_leg( request\_id INTEGER, handler\_id INTEGER, start\_time TIMESTAMP, destination\_facility INTEGER )
- Insert a leg for the delivery request whose ID equals request\_id in the input parameters
  - The function should set the leg\_id of the leg properly
  - In the input parameters, handler\_id refers to handler\_id of the leg; similarly for start\_time and destination\_facility
  - The end\_time of the leg should be set to NULL

## 2.2 Functions

The following routines in this section have return values, and should be implemented as PostgreSQL functions (<https://www.postgresql.org/docs/14/sql-createfunction.html>).

- (1) view\_trajectory( request\_id INTEGER )
- Output: TABLE ( source\_addr TEXT, destination\_addr TEXT, start\_time TIMESTAMP, end\_time TIMESTAMP )
  - Retrieves the legs and return\_legs (if any) of a delivery request, and outputs the addresses of the starting point and ending point of each leg and return\_leg, as well as their start\_time and end\_time, in ascending order of start\_time
  - request\_id refers to the ID of the delivery request
- (2) get\_top\_delivery\_persons ( k INTEGER )
- Output: TABLE ( employee\_id INTEGER )
  - Retrieves the IDs of the top k delivery persons that have handled the largest numbers of “trips”, in descending order of the numbers of trips
    - A trip is defined as a leg, or return delivery leg, or an unsuccessful pickup
    - If two delivery persons have the same number of trips, then the one with the smaller id should rank before the other

(3) `get_top_connections( k INTEGER )`

- Output: TABLE( `source_facility_id` INTEGER, `destination_facility_id` INTEGER )
- Finds the top k “connections” that have the largest numbers of total occurrences in the legs and return\_legs tables
  - A connection is defined as a pair of `source_facility` and `destination_facility`
  - Each time a leg or a return leg starts from a `source_facility` S and ends at a `destination_facility` T, we say that the connection (S, T) “occurs” once
  - In the output, you should rank the connections in descending order of the numbers of occurrences; if two connections have the same number of total occurrences, then they should be sorted in ascending order of (`source_facility_id`, `destination_facility_id`)

### 3. Deadline and Deliverables

Each team is to upload a zip file named **teamNNN.zip**, where NNN is the 3-digits team number, to the Canvas assignment named “**Triggers and Routines**”. You should add leading zeroes to your team number (e.g., team005.zip). Only one file is to be submitted per group. If there are multiple submissions for a group, only the latest submission will be chosen. If the latest submission is uploaded after the deadline, the penalty applies.

The submission **teamNNN.zip** should contain the following two files:

- report.pdf: project report in pdf format
- proc.sql: the triggers and routines of your implementation

The project report (up to a maximum of 20 pages in pdf format with at least 10-point font size) should include the following contents:

- Names and student numbers of all team members and project team number (on the first page).
- A listing of the Project (Part 2) responsibilities of each team member.
  - Team members who did not contribute a fair share of the project work will not receive the same marks awarded to the team.
- For each trigger:
  - Provide the name of the trigger.
  - Explain the basic idea of the trigger implementation.
- For each routine:
  - Explain the basic idea of the routine implementation.
- A summary of any difficulties encountered and lessons learned from the project.

For late submissions, 2 marks (out of 10) will be deducted for submissions up to one day’s late; submissions late for more than one day will receive zero marks and will not be graded.

### 4. FAQ

(1) “Why are we not allowed to make any changes to the schema provided or the routine names and parameters given?”

- This is because we will evaluate each submission using automatic testing. That is, we will prepare some test data following the schema provided to you, and then combine it with your proc.sql and run some queries. We will examine the database state after each query to check whether it is correct; if it is, then you get some marks associated with the query. To facilitate such automatic testing, it is important that every team’s implementation is based on exactly the same schema, and uses the exact routine interface that we provided.

(2) "Will the test data be provided?"

- Unfortunately, no. But you may generate your own test data to check your implementation.